

Doing Data Science in R

Garrick Aden-Buie // April 11, 2014

INFORMS Code & Data Boot Camp



In this project we'll do a simple data science project based on the **Kaggle Titanic Challenge**.

Overview

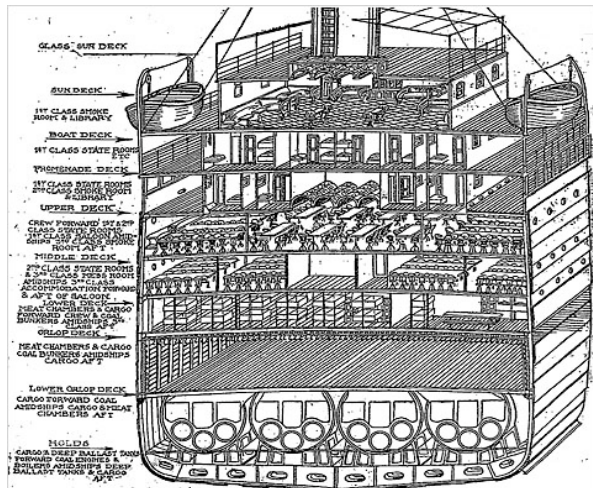
- ▶ Data Exploration
- ▶ Data Cleaning
- ▶ Training a Model
- ▶ Fitting a Model

Disclaimer: Draws heavily from

<http://statsguys.wordpress.com/2014/01/03/first-post/> and <https://github.com/wehrley/wehrley.github.io/blob/master/SOUPUTONUTS.md>.



Who survives the Titanic?



Getting started

- ▶ Download the CSV and R script file from <http://bit.ly/USFCodeCamp2014>
- ▶ Open the R script
- ▶ Set your working directory



The data



Loading the data

```
titanic <- read.csv('titanic.csv', header = TRUE,  
                  na.strings=c('NA', ''))  
  
titanic$Survived <- factor(titanic$Survived,  
                          labels=c('No', 'Yes'))  
  
titanic$Pclass <- factor(titanic$Pclass)
```



Quick look at the data

```
names(titanic)
```

```
## [1] "PassengerId" "Survived"    "Pclass"      "Name"        "Sex"  
## [6] "Age"          "SibSp"       "Parch"       "Ticket"      "Fare"  
## [11] "Cabin"        "Embarked"
```

Also look at:

```
head(titanic)
```

```
summary(titanic)
```

```
str(titanic)
```




Variable Meanings

Variable	Meaning
survival	Survival (0 = No; 1 = Yes)
pclass	Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)
name	Name
sex	Sex
age	Age
sibsp	Number of Siblings/Spouses Aboard
parch	Number of Parents/Children Aboard
ticket	Ticket Number
fare	Passenger Fare
cabin	Cabin
embarked	Port of Embarkation (C = Cherbourg; Q = Queenstown; S = Southampton)



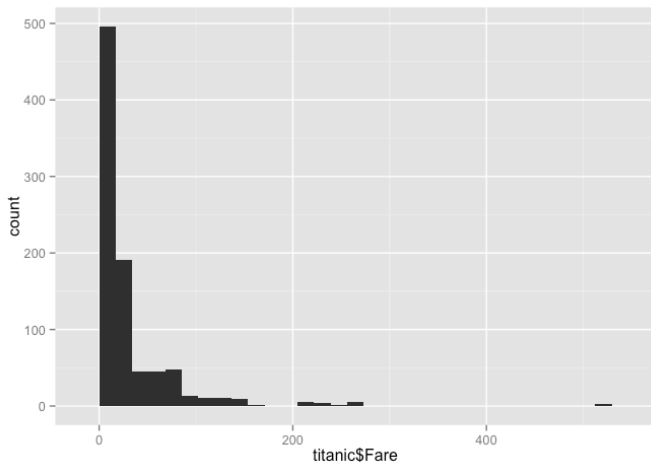
Plotting age

```
require(ggplot2)
qplot(titanic$Age, geom='histogram')
```



Plot Fare

```
qplot(titanic$Fare, geom='histogram')
```



Survival by gender

```
table(titanic$Survived, titanic$Sex)
```

```
##
```

```
##      female male
```


```
## No         81  468
```

```
## Yes        233  109
```



Survival by gender plot

```
ggplot(titanic, aes(x=Sex, fill=Survived))+geom_histogram()
```



Survival by Passenger Class

```
table(titanic$Survived, titanic$Pclass)
```

```
##
```

```
##           1    2    3
```

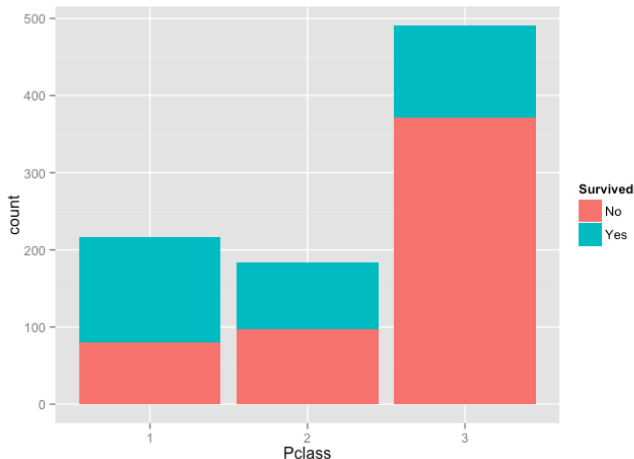
```
##    No    80   97  372
```

```
##    Yes  136   87  119
```




Survival by Passenger Class plot

```
ggplot(titanic, aes(x=Pclass, fill=Survived))+  
  geom_histogram(binwidth=1)
```



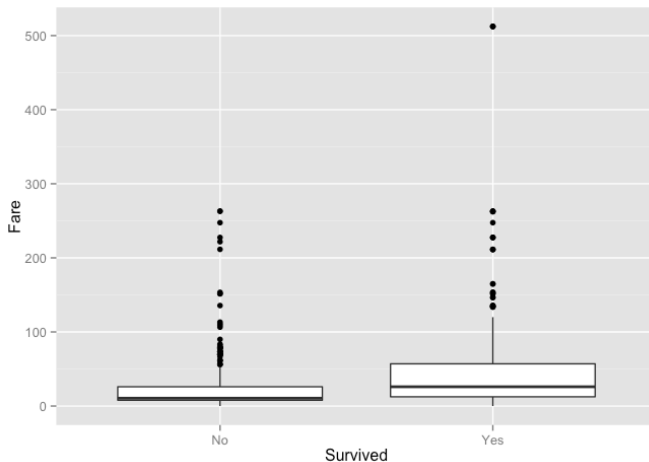
Survival by Age

```
ggplot(titanic, aes(x=Survived, y=Age))+geom_boxplot()
```



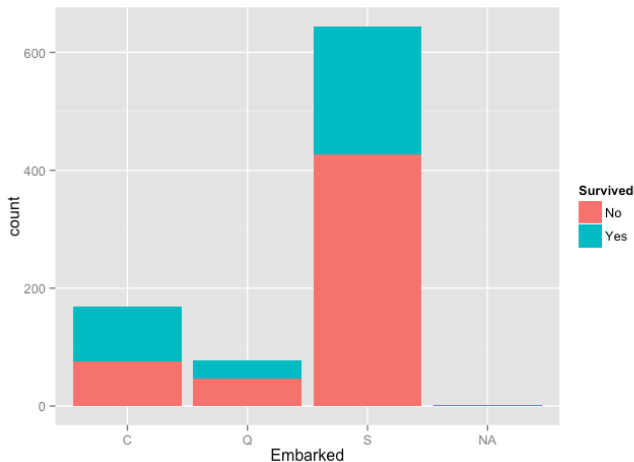
Survival by Fare

```
ggplot(titanic, aes(x=Survived, y=Fare))+geom_boxplot()
```



Survival by Port

```
ggplot(titanic, aes(x=Embarked, fill=Survived))+geom_histogram()
```




Thoughts?



Thoughts?

```
require(Amelia)
missmap(titanic, col=c('blue', 'grey'))
```



Cleaning the data



Missing values

Clearly we need to work on the missing values. Let's ignore Cabin and drop passengers missing Embarked.

```
names(titanic)
```

```
## [1] "PassengerId" "Survived" "Pclass" "Name" "Sex"  
## [6] "Age" "SibSp" "Parch" "Ticket" "Fare"  
## [11] "Cabin" "Embarked"
```

```
titanic <- titanic[, -11]  
titanic <- titanic[!is.na(titanic$Embarked),]
```

But we definitely need to fix Age

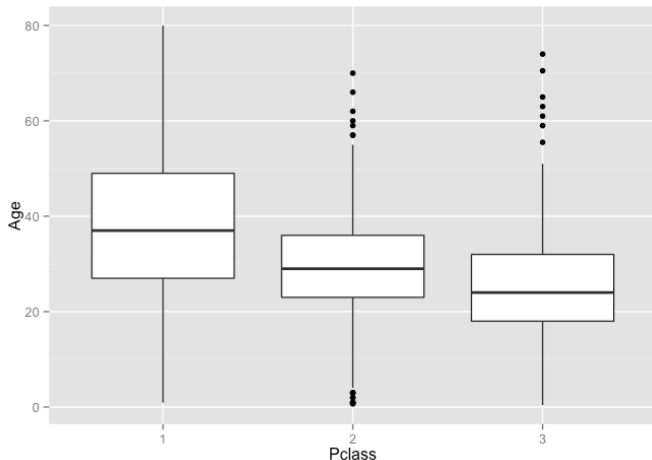
```
length(titanic[is.na(titanic$Age), 'Age'])/dim(titanic)[1]
```

```
## [1] 0.199
```



Does Passenger Class Help?

```
ggplot(titanic, aes(x=Pclass, y=Age))+geom_boxplot()
```



What about the passenger names?

```
rrows <- c(766, 490, 509, 384, 34,  
          126, 887, 815, 856, 851)  
titanic[rrows, 'Name']
```

```
## [1] Brewster, Dr. Arthur Jackson  
## [2] Hagland, Mr. Konrad Mathias Reiersen  
## [3] Lang, Mr. Fang  
## [4] Plotcharsky, Mr. Vasil  
## [5] Wheadon, Mr. Edward H  
## [6] McMahon, Mr. Martin  
## [7] Johnston, Miss. Catherine Helen "Carrie"  
## [8] Fry, Mr. Richard  
## [9] Daly, Mr. Peter Denis  
## [10] Boulos, Miss. Nourelain  
## 891 Levels: Abbing, Mr. Anthony ... Zimmerman, Mr. Leo
```



Passenger Titles

The following titles have at least one person missing Age

- ▶ Dr.
- ▶ Master.
- ▶ Miss.
- ▶ Mr.
- ▶ Mrs.

These titles are clearly correlated with passenger age.



How we're going to do this

- Find indexes of Names that contain Dr.

```
dr <- grep('Dr.', titanic$Name, fixed=TRUE); dr
```

```
## [1] 245 317 398 632 660 766 796
```

- Calculate median age for those passengers

```
m_age <- median(titanic[dr, 'Age'], na.rm=TRUE); m_age
```

```
## [1] 46.5
```

- Select indexes that are both missing and have Dr.

```
dr[dr %in% which(is.na(titanic$Age))]
```

```
## [1] 766
```



Impute Age with median age for titles

```
titles <- c('Dr.', 'Master.', 'Miss.', 'Mr.', 'Mrs.')

for(title in titles){
  passengers <- grep(title, titanic$Name, fixed=TRUE)
  median_age <- median(titanic[passengers, 'Age'], na.rm=TRUE)
  titanic[passengers[passengers %in% which(is.na(titanic$Age))],
    'Age'] <- median_age
}
```



Adding features: Child?

Add a feature to indicate if the passenger is a child (<12)

```
titanic$Child <- 'No'  
titanic[titanic$Age <= 12, 'Child'] <- 'Yes'  
titanic$Child <- factor(titanic$Child)  
summary(titanic$Child)
```

```
##  No  Yes
```

```
## 816  73
```



Adding features: Mother?

Add a feature to indicate if the passenger is a mother. Use the variable `Parch` and title `'Mrs.'`



Adding features: Mother?

Add a feature to indicate if the passenger is a mother. Use the variable Parch and title 'Mrs.'

```
titanic$Mother <- 'No'
mrs <- grep('Mrs.', titanic$Name, fixed=TRUE)
parent <- which(titanic$Parch > 0)
titanic[mrs %in% parent, 'Mother'] <- 'Yes'
titanic$Mother <- factor(titanic$Mother)
summary(titanic$Mother)
```

```
##  No Yes
```

```
## 493 396
```



Divide the data



Divide the data into training and testing sets.

We'll use the caret package for this.

```
require(caret)  
require(pROC)  
require(e1071)
```

<http://caret.r-forge.r-project.org/>

Can be used as a power tool to test and train models.



Make a training and testing set

```
train_index <- createDataPartition(y=titanic$Survived,  
                                     p=0.80,  
                                     list=FALSE)
```

```
train <- titanic[ train_index,]  
test  <- titanic[-train_index,]
```

```
dim(train)
```

```
## [1] 712  13
```

```
dim(test)
```

```
## [1] 177  13
```



Build some models!



Generalized Linear Model (logistic regression)

```
train.glm <- glm(Survived ~ Pclass + Sex + Age +  
                  Child + Sex+Pclass + Mother +  
                  Embarked + Fare,  
                  family = binomial,  
                  data = train)
```



Model summary

```
train.glm
```

```
##  
## Call: glm(formula = Survived ~ Pclass + Sex + Age + Child + Sex + Pclass +  
##      Mother + Embarked + Fare, family = binomial, data = train)  
##  
## Coefficients:  
## (Intercept)      Pclass2      Pclass3      Sexmale      Age      ChildYes  
##      3.182712     -0.593098     -2.127109     -2.506753     -0.020691     0.965256  
##      MotherYes      EmbarkedQ      EmbarkedS      Fare  
##      0.093608      0.058169     -0.671614      0.000492  
##  
## Degrees of Freedom: 711 Total (i.e. Null);  702 Residual  
## Null Deviance:      947  
## Residual Deviance: 637  AIC: 657
```



```
anova(train.glm, test='Chisq')
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Survived
##
## Terms added sequentially (first to last)
##
##
```

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
## NULL			711	947	
## Pclass	2	86.8	709	860	< 2e-16 ***
## Sex	1	188.9	708	671	< 2e-16 ***
## Age	1	19.6	707	652	9.5e-06 ***
## Child	1	4.9	706	647	0.026 *
## Mother	1	0.1	705	647	0.760
## Embarked	2	9.7	703	637	0.008 **
## Fare	1	0.1	702	637	0.822

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Set up caret to train models for us

This just reduces repeated typing later

```
cv.ctrl <- trainControl(method = 'repeatedcv',  
                        repeats = 3,  
                        summaryFunction = twoClassSummary,  
                        classProbs = TRUE)
```



Train glm with caret

```
glm.train <- train(Survived ~ Pclass + Sex +  
  Age + Child + Embarked,  
  data = train,  
  method = 'glm',  
  metric = 'ROC',  
  trControl = cv.ctrl)
```



Check results

```
glm.train
```

```
## Generalized Linear Model
##
## 712 samples
## 12 predictors
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 3 times)
##
## Summary of sample sizes: 641, 641, 640, 641, 640, 641, ...
##
## Resampling results
##
##   ROC   Sens   Spec   ROC SD   Sens SD   Spec SD
##   0.8   0.9    0.7    0.05    0.06     0.08
##
##
```



More details

```
summary(glm.train)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.671   -0.728   -0.362    0.641    2.475
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  3.27714    0.48035   6.82  9.0e-12 ***
## Pclass2     -0.62824    0.29361  -2.14   0.032 *
## Pclass3     -2.17345    0.28058  -7.75  9.5e-15 ***
## Sexmale     -2.51053    0.21020 -11.94 < 2e-16 ***
## Age         -0.02083    0.00975  -2.14   0.033 *
## ChildYes     0.97452    0.41203   2.37   0.018 *
## EmbarkedQ    0.05882    0.40999   0.14   0.886
## EmbarkedS   -0.67022    0.26379  -2.54   0.011 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 947.02  on 711  degrees of freedom
## Residual deviance: 637.24  on 704  degrees of freedom
## AIC: 653.2
##
## Number of Fisher Scoring iterations: 5
```



Random forest model

Let's try the method known as *random forests*.

```
set.seed(42)
rf.train <- train(Survived ~ Pclass + Sex +
                  Age + Child + Embarked,
                  data = train,
                  method = 'rf',
                  metric = 'ROC',
                  trControl = cv.ctrl)
```



Random forests results

```
rf.train
```

```
## Random Forest
##
## 712 samples
## 12 predictors
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 3 times)
##
## Summary of sample sizes: 641, 641, 641, 641, 641, 641, ...
##
## Resampling results across tuning parameters:
##
## mtry ROC Sens Spec ROC SD Sens SD Spec SD
## 2 0.9 1 0.6 0.04 0.03 0.09
## 4 0.9 0.9 0.6 0.05 0.03 0.09
## 7 0.8 0.9 0.7 0.06 0.05 0.1
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```



Compare performance



Make our predictions

```
glm.pred <- predict(glm.train, test)
rf.pred  <- predict(rf.train, test)

glm.prob <- predict(glm.train, test, type='prob')
rf.prob  <- predict(rf.train, test, type='prob')
```



glm prediction results

```
confusionMatrix(glm.pred, test$Survived)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction No Yes
##           No  95  24
##           Yes 14  44
##
##           Accuracy : 0.785
##           95% CI : (0.717, 0.843)
##           No Information Rate : 0.616
##           P-Value [Acc > NIR] : 1.07e-06
##
##           Kappa : 0.533
##           McNemar's Test P-Value : 0.144
##
##           Sensitivity : 0.872
##           Specificity : 0.647
##           Pos Pred Value : 0.798
##           Neg Pred Value : 0.759
##           Prevalence : 0.616
##           Detection Rate : 0.537
##           Detection Prevalence : 0.672
##           Balanced Accuracy : 0.759
##
##           'Positive' Class : No
##
```



randomForest results

```
confusionMatrix(rf.pred, test$Survived)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No Yes
##      No   102  26
##      Yes    7   42
##
##           Accuracy : 0.814
##           95% CI : (0.748, 0.868)
##      No Information Rate : 0.616
##      P-Value [Acc > NIR] : 1.06e-08
##
##           Kappa : 0.584
##  McNemar's Test P-Value : 0.00173
##
##           Sensitivity : 0.936
##           Specificity : 0.618
##           Pos Pred Value : 0.797
##           Neg Pred Value : 0.857
##           Prevalence : 0.616
##           Detection Rate : 0.576
##      Detection Prevalence : 0.723
##           Balanced Accuracy : 0.777
##
##           'Positive' Class : No
##
```




pROC objects for ROC curves

```
glm.ROC <- roc(response = test$Survived,  
               predictor = glm.prob$Yes,  
               levels = levels(test$Survived))  
  
rf.ROC  <- roc(response = test$Survived,  
               predictor = rf.prob$Yes,  
               levels = levels(test$Survived))
```



ROC Plot

```
plot(glm.ROC)  
plot(rf.ROC, add=TRUE, col="red")
```



Thanks!

